

Exercise 1

Simple data processing with Python and Pandas

Prior Knowledge

Unix Command Line Shell

Simple Python

Learning Objectives

First steps with Pandas

Understand the Jupyter Notebook model

Software Requirements

(see separate document for installation of these)

- Python 2.7.14
- Jupyter notebooks
- (This exercise should also work using Python 3.6)

Ensuring we have the pre-requisite code installed.

1. Before we can use *pandas* we need to make sure we have it installed.
From the VM's command-line (terminal window), enter:

```
sudo pip install pandas
sudo pip install matplotlib
```

This should ensure that *pandas*, *matplotlib* and any dependencies are downloaded and installed on your VM. **pip** is a Python package manager that installs Python packages from a central source of packages on the Internet. (**P**ip **I**nstalls **P**ackages)

Downloading our sample data

2. Let's make a directory to store our code.

```
mkdir hyg
cd hyg
```

3. Now let's download some star data.

This data is found at:

<http://www.astronexus.com/hyg>

You can either download the data by going to that website and finding HYG3.0 and downloading into the newly created directory, or you can use a command line and type:

```
wget http://www.astronexus.com/files/downloads/hygdata_v3.csv.gz
```

4. Either way that you downloaded it, you now need to uncompress it:

```
gunzip hygdata_v3.csv.gz
```

5. Check it's the right size:

```
ls -l ~/hyg
```

You should see:

```
-rw-rw-r-- 1 big big 33449663 Apr 21 2015 hygdata_v3.csv
```

6. The VM has a “notebook” system called Jupyter configured by default. The result is that instead of starting a command line repl¹, there is a web based editor/evaluator launched instead.

7. To start this, type (from the same command line that is in the hyg directory):

```
jupyter notebook
```

8. In the command-line you will see

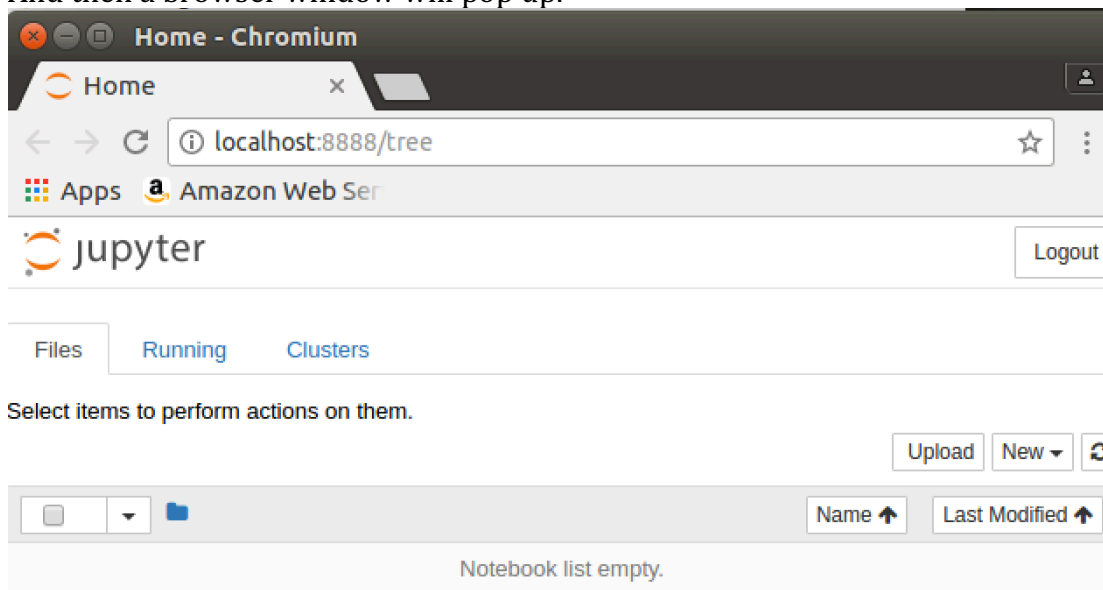
```
[I 13:53:23.865 NotebookApp] Serving notebooks from local
directory: /home/oxclo/pse
[I 13:53:23.866 NotebookApp] 0 active kernels
[I 13:53:23.866 NotebookApp] The Jupyter Notebook is running at:
http://localhost:8888/?token=fd655aab32ed4840ceb47b8b7392b1243a27f5
6350888a91
[I 13:53:23.866 NotebookApp] Use Control-C to stop this server and
shut down all kernels (twice to skip confirmation).
[C 13:53:23.868 NotebookApp]
```

Copy/paste this URL into your browser when you connect for the first time,
to login with a token:

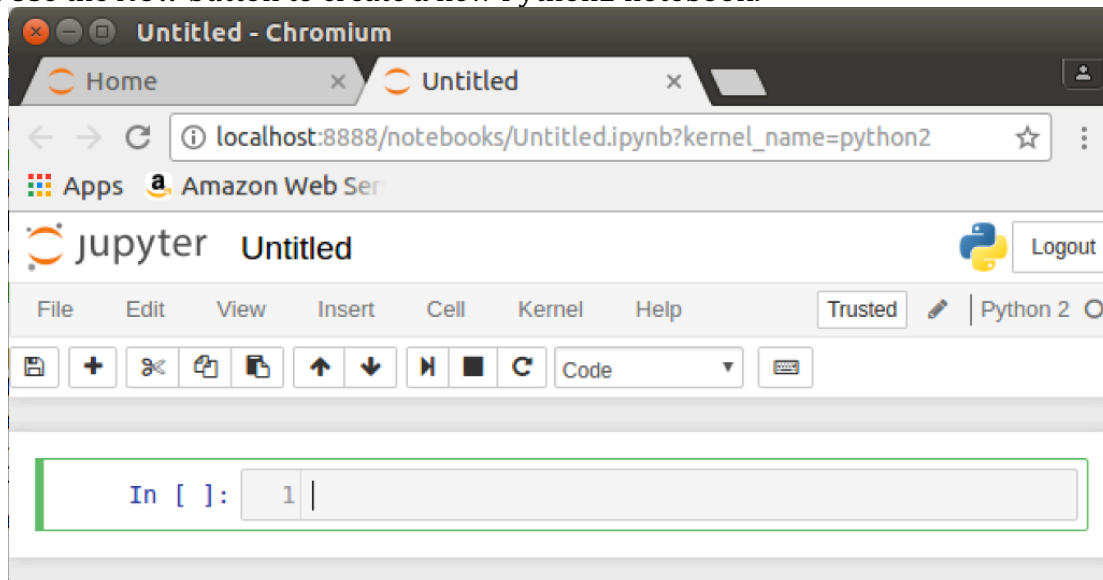
```
http://localhost:8888/?token=fd655aab32ed4840ceb47b8b7392b1243a27f5
6350888a91
```

¹ Read Eval Print Loop

9. And then a browser window will pop up.



10. Use the **New** button to create a new Python2 notebook:




11. Click on the name of the notebook (currently “Untitled”) and rename it to **hyg**
12. Now type the following into the **Cell** (next to the words **In []:**)
You don’t need to type in the comments but it will help you to remember what each bit does!

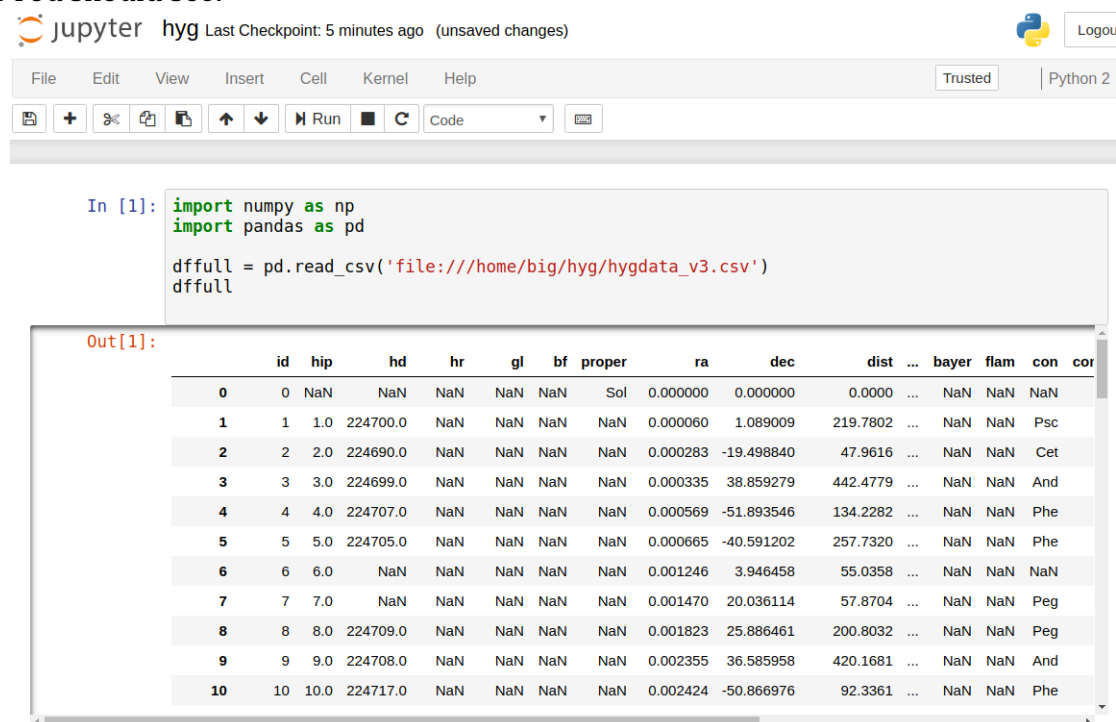
```
import numpy as np # numpy is a library of numerical routines
import pandas as pd # pandas is the data handling library
import matplotlib.pyplot as plt #plotting library

dffull = pd.read_csv('file:///home/big/hyg/hygdata_v3.csv')
dffull # show the
```

This is creating a DataFrame. This is an object offered by the pandas library that helps deal with tabular data. It is very good at dealing with data that naturally falls into rows and columns and also that has missing elements.

13. Now click on the Run icon 

14. You should see:



The Jupyter Notebook interface shows the following code in the input cell:

```
In [1]: import numpy as np
import pandas as pd

dffull = pd.read_csv('file:///home/big/hyg/hygdata_v3.csv')
dffull
```

The output cell displays a DataFrame with 11 rows and 15 columns. The columns are: id, hip, hd, hr, gl, bf, proper, ra, dec, dist, ..., bayer, flam, con, cor. The first 10 rows are visible, with the 11th row partially cut off.

	id	hip	hd	hr	gl	bf	proper	ra	dec	dist	...	bayer	flam	con	cor
0	0	NaN	NaN	NaN	NaN	NaN	Sol	0.000000	0.000000	0.0000	...	NaN	NaN	NaN	
1	1	1.0	224700.0	NaN	NaN	NaN	NaN	0.000060	1.089009	219.7802	...	NaN	NaN	Psc	
2	2	2.0	224690.0	NaN	NaN	NaN	NaN	0.000283	-19.498840	47.9616	...	NaN	NaN	Cet	
3	3	3.0	224699.0	NaN	NaN	NaN	NaN	0.000335	38.859279	442.4779	...	NaN	NaN	And	
4	4	4.0	224707.0	NaN	NaN	NaN	NaN	0.000569	-51.893546	134.2282	...	NaN	NaN	Phe	
5	5	5.0	224705.0	NaN	NaN	NaN	NaN	0.000665	-40.591202	257.7320	...	NaN	NaN	Phe	
6	6	6.0	NaN	NaN	NaN	NaN	NaN	0.001246	3.946458	55.0358	...	NaN	NaN	NaN	
7	7	7.0	NaN	NaN	NaN	NaN	NaN	0.001470	20.036114	57.8704	...	NaN	NaN	Peg	
8	8	8.0	224709.0	NaN	NaN	NaN	NaN	0.001823	25.886461	200.8032	...	NaN	NaN	Peg	
9	9	9.0	224708.0	NaN	NaN	NaN	NaN	0.002355	36.585958	420.1681	...	NaN	NaN	And	
10	10	10.0	224717.0	NaN	NaN	NaN	NaN	0.002424	-50.866976	92.3361	...	NaN	NaN	Phe	

15. Scroll down to the bottom of the table and you should see how many rows (stars) are in the catalogue. Note how the notebook automatically knows how to display pandas dataframes in an intelligent manner. Also note that you are not seeing all the rows or columns because there is too much data to display.

You can see the description of the columns here:

<https://github.com/astronexus/HYG-Database/blob/master/README.md>

16. Before we do any more data processing, let's configure Jupyter to do nice *tab completion*. In a new cell enter:

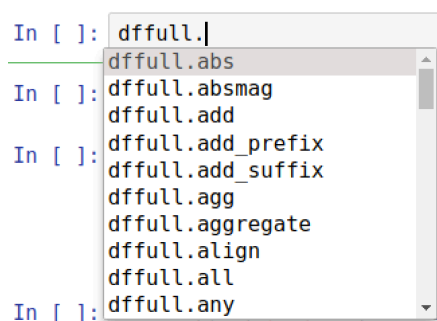
```
%config IPCompleter.greedy=True
```

Anything starting with % is a hint that this is for Jupyter not for Python.

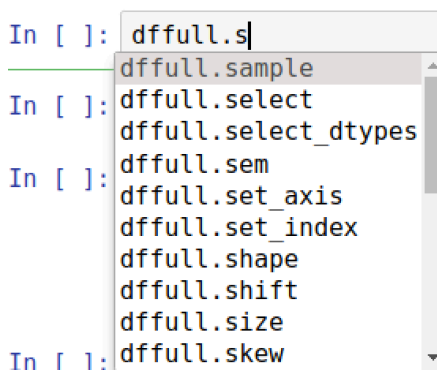
17. You can also just get that information (number of rows and columns) by using the dataframe shape. In the next cell type:

```
dffull.
```

Before typing anything else, hit the Tab key. You should see all possible options for syntax now appear in a little box like this:




Now type 's', and you should see just the operations starting with 's' appear:



Now move down using the down arrow and select 'shape' by hitting Enter.

Now hit **Ctrl-Enter** (same as Run icon)

You should see:



Not all the columns are of interest to us. One simple approach is to create a new dataframe that only uses some of the columns from the old dataframe.

To do that, we can use the following syntax:

```
columns = ['id', 'gl', 'mag', 'absmag', 'proper', 'ra', 'dec',
           'dist', 'con', 'ci', 'lum']
df = pd.DataFrame(dffull, columns=columns)
df # show the resulting dataframe
```

18. Paste or type that into a new cell and execute it.

19. Now, let's identify the stars that have a 'proper' name.

20. There are a couple of ways we could do this. The first, simple one, is to just select that column, and then drop all NaN entries:

Execute:

```
df['proper'].dropna()
```

21. You should see something like:

```
In [21]: df['proper'].dropna()
```

```
Out[21]: 0          Sol
        676      Alpheratz
        744          Caph
        1065      Algenib
        2076      Ankaa
        3172      Shedir
        3413      Diphda
        3759      96 G. Psc
        3820  Van Maanen's Star
        4417          Cih
        5436      Mirach
        6672      Ruchbah
        7574      Achernar
        8884      Sheratan
        9618      Almaak
        9861      Hamal
        10800      Mira
        11734      Polaris
        12082      268 G. Cet
        13813      Acamar
        14100      Menkar
        14540      Algol
        15471      82 G. Eri
```

22. Notice that this no longer looks quite the same. This is because this has created a Series object instead of a DataFrame (Each column is effectively a Series, and we've extracted one column).

23. Suppose we want the whole DataFrame (all the columns), but only those with a 'proper' name. We can use a selection function to *locate* the right

rows:

```
df.loc[df['proper'].notnull()]
```

24. You should see:

```
In [22]: df.loc[df['proper'].notnull()]
```

```
Out[22]:
```

	id	gl	mag	absmag	proper	ra	dec	dist	con	ci	lum
0	0	NaN	-26.70	4.850	Sol	0.000000	0.000000	0.0000	NaN	0.656	1.000000e+00
676	676	NaN	2.07	-0.297	Alpheratz	0.139791	29.090432	29.7442	And	-0.038	1.1449886e+02
744	744	Gl 8	2.28	1.155	Caph	0.152887	59.149780	16.7842	Cas	0.380	3.006076e+01
1065	1065	NaN	2.83	-2.567	Algenib	0.220598	15.183596	120.0480	Peg	-0.190	9.264031e+02
2076	2076	NaN	2.40	0.327	Ankaa	0.438056	-42.305981	25.9740	Phe	1.083	6.444660e+01
3172	3172	NaN	2.24	-1.985	Shedir	0.675116	56.537331	69.9790	Cas	1.170	5.420009e+02
3413	3413	Gl 31	2.04	-0.312	Diphda	0.726490	-17.986605	29.5334	Cet	1.019	1.160914e+02
3759	3759	Gl 33	5.74	6.378	96 G. Psc	0.806382	5.280615	7.4549	Psc	0.890	2.447936e-01
3820	3820	Gl 35	12.37	14.222	Van Maanen's Star	0.819416	5.388610	4.2626	NaN	0.554	1.783200e-04
4417	4417	NaN	2.15	-3.981	Cih	0.945143	60.716740	168.3502	Cas	-0.046	3.407219e+03
5436	5436	Gl 53.3	2.07	-1.840	Mirach	1.162194	35.620558	60.5327	And	1.576	4.742420e+02

```
In [ ]: |
```

25. This has selected every row which meets the criteria (i.e. that the column *proper* is not null).

26. You can sort the data based on a column using the following syntax, e.g. to identify the stars by distance.

```
df.sort_values('dist', ascending=False)
```

27. If you just want to see the first 10 rows of a DataFrame you can use:

```
df.head(n=10)
```

28. Use those to identify the five furthest “proper named” stars. What do you think of the data?

29. You can select on multiple criteria at once, e.g.:

```
df.loc[(df['proper'].notnull()) & (df['dist']<100000)].sort_values('dist', ascending=False)
```

30. Identify the Gliese catalog identifier of the three least luminescent stars.

Visualisation

31. We can do some simple graphing of the data in Jupyter very easily. Matplotlib is a simple graphing package for Python. We have already installed it and imported it. This next line tells Jupyter to automatically plot diagrams made by matplotlib directly in the notebook.

```
%matplotlib notebook
```

32. Any pandas dataframe or series is automatically plottable by matplotlib (although you may not get anything useful!).


33. Try it:

```
df.plot()
```

34. For something more useful, let's plot a scatter graph of luminosity vs distance:

```
df.plot.scatter(x='dist', y='lum')
```

35. You should see:



36. Once again, it looks like the data is incorrect and therefore not useful (see the comment in the documentation under the distance attribute).

37. Redo the graph this time filtering out any distance $\geq 100,000$.

38. This still isn't much use. Now try making the scales logarithmic by adding the parameters `logx=True`, `logy=True` to the plot. Is there anything meaningful about the resulting graph?

39. Extension:

Explore the data further using the matplotlib to identify any interesting correlations between the data.

40. Before finishing, close the Jupyter browser windows and then stop the Jupyter server by using Ctrl-C on the window, and then **y**.

