


Big Data Engineering

Other Tools and Libraries

Julie Weeds
March 2019

© Paul Fremontie 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

HDFS in a nutshell



© Paul Fremontie 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

HDFS inspiration

- Google File System
 - Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. 2003. The Google file system. In Proceedings of the nineteenth ACM symposium on Operating systems principles (SOSP '03). ACM, New York, NY, USA, 29-43.

The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung
Google

ABSTRACT

We have designed and implemented the Google File System, a scalable distributed file system for large distributed data-intensive applications. It provides fault tolerance while meeting an aggressive consistency, latency, and reliability high-availability performance to a large number of clients. While sharing some of the same goals as previous distributed file systems, our design has been driven by observations of our application workloads and technological environment, both current and anticipated, that reflect a needed departure from more earlier file system assumptions. This has led us to implement traditional features and explore previously different design options.

The file system has successfully met our storage needs. It is widely deployed within Google on the storage platform.

1. INTRODUCTION

We have designed and implemented the Google File System (GFS) to meet our rapidly growing demand for large data processing needs. GFS shares many of the same goals as previous distributed file systems such as performance, reliability, scalability, and availability. However, its design has been driven by key observations of our application workloads and technological environment, both current and anticipated, that reflect a needed departure from more earlier file system design assumptions. We have implemented traditional features and explored radically different points in the design space.

First, our application workloads are the worst rather than the best. The file system consists of hundreds or even thousands of storage machines built from inexpensive main

© Paul Fremontie 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

HDFS overview

- Good for streaming access to large files, reliability, scale
- Not good for random access, small files
- Blocks of data 64Mb in size (configurable)
- Each block can be replicated across multiple data nodes for High Availability (HA)

© Paul Fremontie 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

HDFS Usage

- Spotify has 1600+ nodes, storing 60+ petabytes of data
 - <https://www.usenix.org/system/files/conference/fast17/fast17-niazi.pdf>
- One of Facebook's largest clusters (based on HDFS) holds more than 100 PB of data, processing more than 60,000 Hive queries a day
 - <https://www.facebook.com/notes/facebook-engineering/under-the-hood-scheduling-mapreduce-jobs-more-efficiently-with-corona/>

© Paul Fremontie 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

HopFS

- HopFS is a drop-in replacement for HDFS, based on HDFS v2.0.4.



Figure 1: System architecture for HDFS and HopFS. For high availability, HDFS requires an Active NameNode (ANN), at least one Standby NameNode (SbNN), at least three Journal Nodes for quorum-based replication of the write ahead log of metadata changes, and at least three ZooKeeper instances for quorum based coordination. HopFS supports multiple stateless namenodes that access the metadata stored in NDB database nodes.


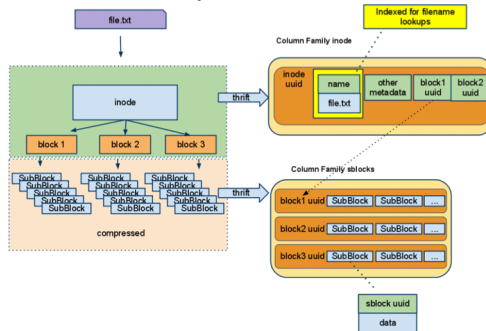


Figure 6: HopFS and HDFS throughput for Spotify workload.

CassandraFS (not open source)



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Amazon S3

- Simple Storage Service
- Unlimited storage of files
 - Up to 5 terabytes each
 - Stored in named “buckets”
 - Accessible via AWS APIs or HTTP
 - Authenticated or Public

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark packages

- A wide set of plugins
 - Currently 148 community donated plugins
- Data connectors
 - Cassandra, Couchbase, Mongo, CSV, etc
- Machine Learning, Neural networks
- Streaming
- etc

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Using Spark Packages

Automatic download from the web:

```
bin/spark-shell
--packages com.databricks:spark-csv_2.11:1.2.0
```


© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Locality

- Spark understands the locality of data:
 - Already in memory
 - HDFS location
 - Cassandra location

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark Extras



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark Extras

- Spark Streaming
 - Realtime analysis in Spark
- Spark MLlib
 - Like Mahout – Machine learning in Spark
- GraphX
 - Graph processing in Spark
- SparkR
 - R statistical analysis on Spark

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark MLlib

- Simple stats and correlation testing
- Classification and regression
- Collaborative Filtering
 - Alternating Least Squares
- Clustering
 - k-means, etc
- Frequent Pattern Mining
- Plus more

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

MLlib example

```
from pyspark.mllib.fpm import FPBGrowth

data = sc.textFile("data/mllib/sample_fpgrowth.txt")

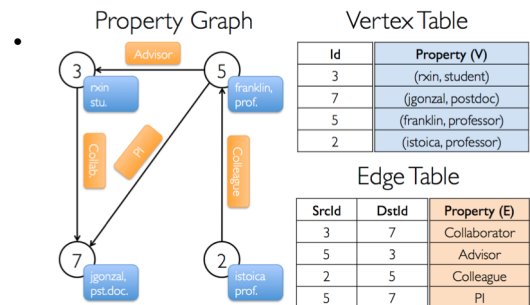
transactions = data.map(lambda line: line.strip().split(' '))

model = FPBGrowth.train(transactions, minSupport=0.2,
                          numPartitions=10)

result = model.freqItemsets().collect()
for fi in result:
    print(fi)
```

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

GraphX



 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

R



- R is an open source system for statistics and graphics
 - Based on the S language from AT&T Bell Labs
- Supports a wide variety of statistical techniques and graphing tools
- An extensible set of packages that provide extra functions via CRAN
 - The Comprehensive R Archive Network

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

SparkR

- A lightweight approach to use Spark from within R
- Also works with MLlib for machine learning
- Allows complex statistical analysis to be done on a Spark cluster

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

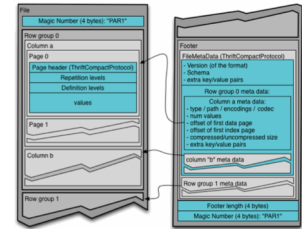
Apache Avro

- A compact data storage and transmission system
 - Uses schemas of data to ensure it can be read by the receiver
 - Supports dynamic typing
- Used by RPC or data collection systems
 - Fast binary protocols
- Also supports storage
 - Hence used by many Big Data apps including Hadoop and Spark

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Apache Parquet

- Apache Parquet is a columnar data storage model
 - Works with Hadoop, Spark and many others
 - Efficient storage of data
 - Based on another Google system called Dremel



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Cluster management systems for Big Data

- YARN
 - Part of Hadoop but significantly rebuilt since Hadoop 1
- Mesos
 - Popular Apache project
 - Built to be a resource manager for a complete datacenter
 - Supports many workloads (e.g. Docker as well as Spark)


© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

What is YARN?

- **YARN is the system that runs your code on multiple nodes**
- Hadoop 2.0 replacement for the cluster manager
 - Basically a model to distribute and manage workloads
 - Not just MapReduce but supports other workloads

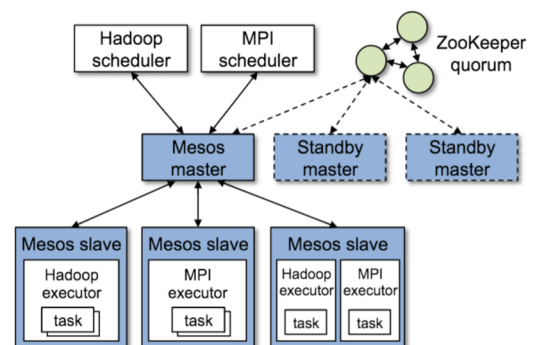
© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

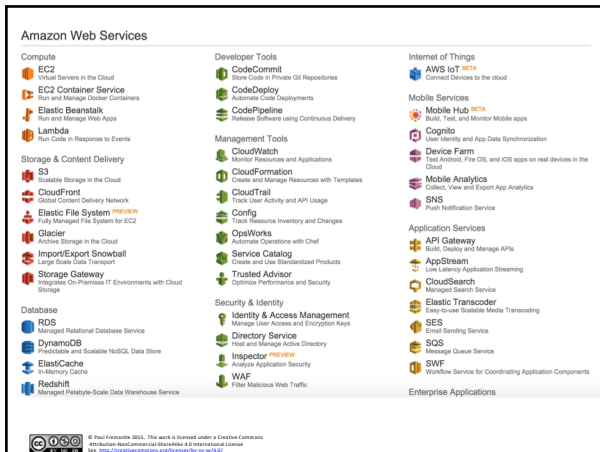
YARN architecture



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Apache Mesos





EC2 / AWS main functions

- EC2 (Elastic Compute Cloud)
 - Instances
 - Servers of various sizes
 - AMIs (Amazon Machine Images)
 - Server images
 - Elastic Block Storage (EBS)
 - Virtualized Hard drives
 - VPC (Virtual Private Cloud)
 - Secure network space
- S3 (Simple Storage Solution)
 - “Buckets” of data
 - Longer term storage of data

Flintrock



license [Apache 2.0](#) build [passing](#) chat [on gitter](#)

Watch [@nchammas's](#) talk on Flintrock at Spark Summit East 2016: [talk](#) / [slides](#)

Flintrock is a command-line tool for launching [Apache Spark](#) clusters.

Installing standalone version

Standalone version (Python not required!)

If you don't have a recent enough version of Python, or if you don't have Python installed at all, you can still use Flintrock. We publish standalone packages of Flintrock on GitHub with our [releases](#).

Find the standalone package for your OS under our [latest release](#), unzip it to a location of your choice, and run the `flintrock` executable inside.

For example:

```
flintrock_version="0.11.0"

curl --location --remote-name "https://github.com/nchammas/flintrock/releases/download/v$flintrock_version/flintrock-$flintrock_version-standalone-OSX-x86_64.zip"
cd flintrock/

# You're good to go!
./flintrock --help
```

You'll probably want to add the location of the Flintrock executable to your `PATH` so that you can invoke it from any directory.

Flintrock Launching a cluster

```
flintrock launch test-cluster \
  --num-slaves 1 \
  --spark-version 2.2.0 \
  --ec2-key-name key_name \
  --ec2-identity-file /path/to/key.pem \
  --ec2-ami ami-a4c7edb2 \
  --ec2-user ec2-user
```

Flintrock configure

```
config.yaml
services:
  spark:
    version: 2.2.0
    hdfs:
      version: 2.7.3
  provider: ec2
providers:
  ec2:
    key-name: bigkp # for west-2
    identity-file: /home/bigkp/keys/bigkp.pem
    instance-type: t2.micro
    region: eu-west-2
    ami: ami-01419b84382864e4 # Amazon Linux, eu-west-2
    user: ec2-user
    tenancy: default # default | dedicated
    ebs-optimized: no # yes | no
    instance-initiated-shutdown-behavior: terminate # terminate | stop

launch:
  num-slaves: 2
  install-hdfs: False
```

Other things you can do

```
flintrock destroy test-cluster
flintrock login test-cluster
flintrock describe test-cluster
flintrock add-slaves test-cluster
    --num-slaves 2
flintrock remove-slaves test-cluster
    --num-slaves 1
flintrock run-command test-cluster
'sudo yum install -y package'
flintrock copy-file test-cluster
/local/path /remote/path
```

 © Paul Foreman 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Questions?

 © Paul Foreman 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>