

Big Data Engineering Realtime Big Data Processing

Julie Weeds
March 2019

© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Streaming

- Continuous data flow
 - “Unbounded streams of data”
- Usually uses a message distribution system
 - JMS, Apache Kafka, ZeroMQ, MQTT
- An unbounded set of events with time
 - $\langle t_1, E_1 \rangle, \langle t_2, E_2 \rangle, \dots, \langle t_n, E_n \rangle, \dots$

© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>


Stream processing categorization

- Simple event processing
 - Working on an event at a time
 - e.g. filter out all events where the wind speed > 50 mph
- Event stream processing
 - Time-based processing of a single stream of events
 - Average wind speed over the last hour compared to the average over the last day
- Complex Event Processing
 - Correlation of events across different streams
 - Emergency calls correlated with wind speed in real time

© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Comparing Databases with Real-Time systems

	Database Applications	Event-driven Applications
Query Paradigm	Ad-hoc queries or requests	Continuous standing queries
Latency	Seconds, hours, days	Milliseconds or less
Data Rate	Hundreds of events/sec	Tens of thousands of events/sec or more




© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Approaches to Streaming

- Pure streaming
 - Each event is processed as it comes in
- Micro-batch
 - Small batches of events are processed
 - Typically trades flexibility for performance
- Shared nothing
 - You can process events on any system in the cluster
- Stateful / Partitioned
 - The event must be processed on a system that has the correct state in memory

© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Lambda Architecture



© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution Non-Commercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Data distribution

- You need to get the events to the processing systems


© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

MQTT (Message Queuing Telemetry Transport)

- Very simple, lightweight, fast
- No built in support for clustering / big-data
 - But can make up for it by being very fast
- Used a lot in IoT

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Apache Kafka



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Source: <http://www.slideshare.net/charmalo/>

Kafka

- Many of the approaches we've seen:
 - Partitioning
 - Multiple brokers
 - Elastically scalable
 - Supports clusters of co-ordinated consumers
 - Automatic re-election of leaders

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Kafka exactly-once semantics



Mathias Verraes

@mathiasverraes

Follow


There are only two hard problems in distributed systems: 2. Exactly-once delivery 1. Guaranteed order of messages 2. Exactly-once delivery

RETWEETS

6,775

LIKES

4,727




10:40 AM - 14 Aug 2015

69 6.8K 4.7K

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

ZeroMQ




© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Processing the data

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Apache Storm



Note: another DAG

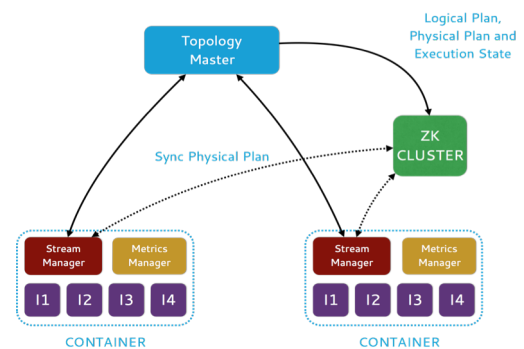
© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Apache Storm

- Originally developed by BackType
 - Nathan Marz
- Acquired by Twitter
- Open Sourced and then donated to Apache
- Became a top level project in 2014
 - <http://storm.apache.org>

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Heron



© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Heron: Key Features

- Fully API compatible with Apache Storm
- Task isolation
- Developer productivity
- Ease of manageability
- Use of mainstream languages C++/Java/Python

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

17

Heron

- In production at Twitter for >2 years




Fig. 2. Throughput with acks


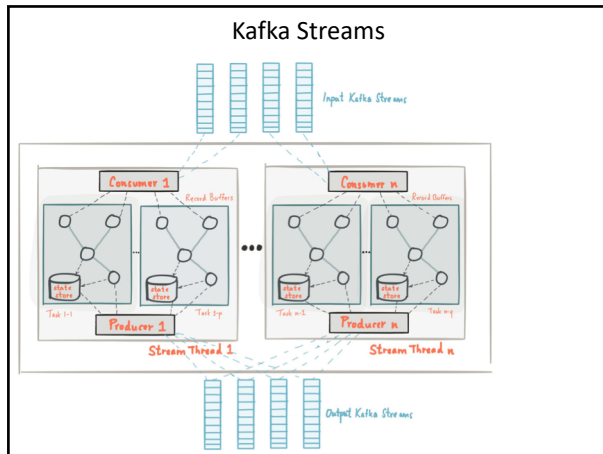


Fig. 3. End-to-end latency with acks

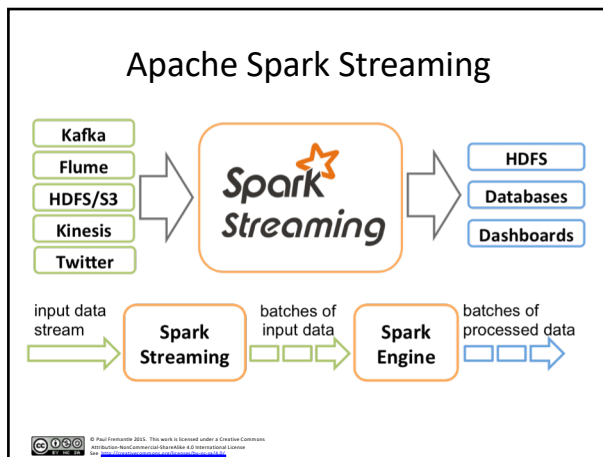
© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>



Kafka Streams


- Event-at-a-time processing (not microbatch) with millisecond latency
- Stateful processing including distributed joins and aggregations
- A convenient DSL
- Windowing with out-of-order data using a DataFlow-like model
- Distributed processing and fault-tolerance with fast failover
- Reprocessing capabilities so you can recalculate output when your code changes
- No-downtime rolling deployments

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

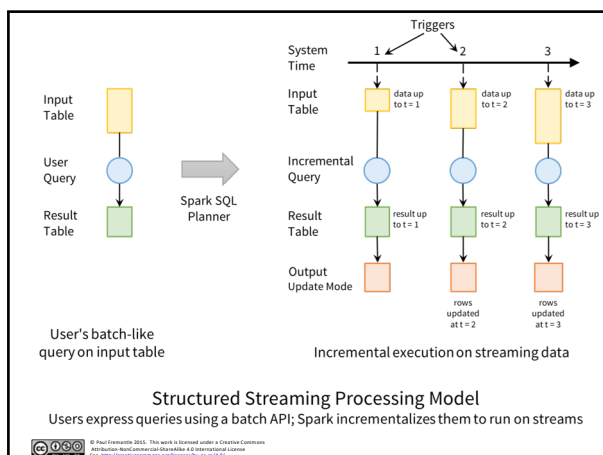


Structured Streams in Spark

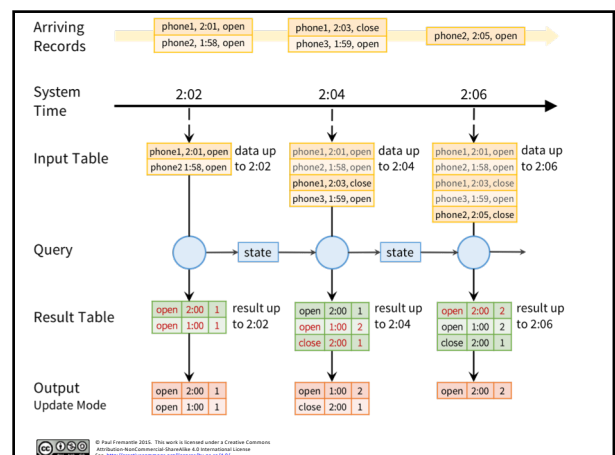
- Since Spark 2.0, there is a much better approach



© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>



© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>



© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>


Siddhi

- A stateful query model
- SQL-like language for querying streams of data
 - Extended with **windows**
 - Time, Event count, batches
 - Partitioned
 - Based on data in the events
 - Pattern matching
 - A then B then C within window

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Siddhi

- Apache Licensed Open Source on Github
 - <https://github.com/wso2/siddhi/>
- Pluggable into Storm, Spark and Kafka Streams
- Supports millions of events/sec
- <http://freo.me/DEBS Siddhi>




SiddhiQL

```
FROM login_stream#window.time(10 min)
SELECT ip,
       count(ip) as loginCount,
       cityId
GROUP BY ip
HAVING loginCount > 10
INSERT INTO login_attemp_repeatedly_stream;
```

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Siddhi at Uber



Siddhi at Uber

- 100+ production apps
- 30 billion messages / day
- Fraud, anomaly detection
- Marketing, promotion
- Monitoring, feedback
- Real time analytics and visualization

<https://freo.me/siddhi-uber>

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Summary

- Realtime processing is hard
 - Requires large memory and state
 - The lambda architecture splits the problem into batch and realtime challenges
- Multiple approaches:
 - Pure Streaming
 - Micro-batch
 - CEP

© Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Questions?

 © Paul Fremstad 2015. This work is licensed under a Creative Commons
Attribution NonCommercial ShareAlike 4.0 International License.
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>